

```
In [1]: 1 import pandas as pd
        2 import sklearn
        3 import matplotlib.pyplot as plt
        4 import numpy as np
        5 import os
        6 from PIL import Image
        7 from sklearn.model_selection import train_test_split
        8 import shutil
        9 import glob
       10 import cv2
```

```
In [2]: 1 import tensorflow as tf
        2 # from tf import keras
        3
        4 import keras
        5 from keras import layers
        6 import visualkeras
        7 from keras.preprocessing.image import ImageDataGenerator
        8 from keras import Sequential
        9 from keras.layers import Dropout , Dense, Flatten, Conv2D, MaxPool2D, BatchN
       10 from tensorflow.keras.optimizers import RMSprop, Adam, SGD, Adadelata
       11 from tensorflow.keras.models import load_model
       12 from keras import Sequential
       13 from keras.layers import LeakyReLU
```

## About the Data

The dataset has been taken from kaggle competition "I'm Something of a Painter Myself" which provides Monets and photos with the objective of training on the monet painting and styling them on the photos. The dataset provided has 2 formats of the same data i.e. '.tfrec' and '.jpeg'. In our scenario, we will be sticking to the jpegs for further work.

The dataset has 300 monets and 7038 photos. The model will be trained using DCGANs.

## Importing Data

```
In [3]: 1 monet_rt = "gan-getting-started/monet_jpg/"
        2 photo_rt = "gan-getting-started/photo_jpg/"
```

```
In [4]: 1 monet_data = os.listdir("gan-getting-started/monet_jpg/")
        2 photo_data = os.listdir("gan-getting-started/photo_jpg/")
```

## Exploratory Data Analysis


```
In [49]: 1 len(monet_data)
```

Out[49]: 300


```
In [50]: 1 len(photo_data)
```

Out[50]: 7038

```
In [7]: 1 fig , ax = plt.subplots(1,4, figsize = (18,10))
2 for i in range(4):
3     ax[i].imshow(Image.open(monet_rt + monet_data[i]))
```



```
In [8]: 1 fig , ax = plt.subplots(1,4, figsize = (18,10))
2 for i in range(4):
3     ax[i].imshow(Image.open(photo_rt + photo_data[i]))
```



## Analysis:

A brief analysis of the images doesn't seem to show anything peculiar or other than the ordinary. They are simply monets which are 300 to train and 7038 photos provided. However, there does seem to be one thing that is clear, that they are all regarding sceneries and they are 256 x 256 pixels.

## Data Preprocessing

In order to put our data into the modelling phase, we need to perform certain preliminary steps to get the data in the correct form. To this, we perform the following actions:

1. Read each image from the each directory(Monet and Photos)
2. Read the image
3. Resize the image to 64 x 64 from 256 x 256
4. Convert into array using keras
5. Append the list containing the images converted to arrays

## Data Scaling and Augmentation

```
In [45]: 1 train_photo_art = []
2         for directory in glob.glob('gan-getting-started/photo_jpg/'):
3             for filename in glob.glob(directory + '/*'):
4                 image = cv2.imread(filename)
5                 image = cv2.resize(image, (64, 64))
6                 image = tf.keras.preprocessing.image.img_to_array(image)
7                 train_photo_art.append(image)
```

```
In [30]: 1 train_images_art = []
2         for directory in glob.glob('gan-getting-started/monet_jpg/'):
3             for filename in glob.glob(directory + '/*'):
4                 image = cv2.imread(filename)
5                 image = cv2.resize(image, (64, 64))
6                 image = tf.keras.preprocessing.image.img_to_array(image)
7                 train_images_art.append(image)
8
9
```

```
In [31]: 1 train_images_art = np.array(train_images_art, dtype="float")
2         train_images_art = (train_images_art - 127.5) / 127.5
```

```
In [32]: 1 train_images_art.shape
```

```
Out[32]: (300, 64, 64, 3)
```

```
In [47]: 1 train_photo_art = np.array(train_photo_art, dtype="float")
2         train_photo_art = (train_photo_art - 127.5) / 127.5
3         train_photo_art.shape
```

```
Out[47]: (7038, 64, 64, 3)
```

## Model

We will be designing the model based on Deep Convolutional Generative Adversarial Network which is more commonly known as DCGAN. The model design and execution will be in the following sequence:

1. Image Generator
2. Image Discriminator
3. joining both to form the GAN
4. Generating results

```
In [33]: 1 dim = 100
```

## Generator

We define the image generator to be a sequential model with an initial dense layer which is

configured to have the dimensions as 8 x 8 x256. We then have 3 hidden layers with each using a transposed convolution layer with 128 filters each, size and strides set to 2 and with evenly padding around. The final layer is a conv2D layer configured to give an output of 64 x 64 image. The model is compiled with the loss set to binary\_crossentropy and optimizer set to Adam. The leakyReLU is also incorporated in each layer to avoid a dead gradient by providing a small arbitrary value for the gradient with the arbitrary value set to 0.3.


```
In [34]: 1 def image_generator():
2         gen = Sequential()
3
4         # input layer
5         gen.add(Dense(8 * 8 * 256 , input_shape = (dim,)))
6         gen.add(LeakyReLU(0.3))
7
8         gen.add(Reshape((8,8,256)))
9
10        # First Hidden Layer
11        gen.add(Conv2DTranspose(128 , kernel_size = 2, strides = 2 , padding = '
12        gen.add(LeakyReLU(0.3))
13
14        # Second Hidden Layer
15        gen.add(Conv2DTranspose(128 , kernel_size = 2, strides = 2 , padding = '
16        gen.add(LeakyReLU(0.3))
17
18        # Third Hidden Layer
19        gen.add(Conv2DTranspose(128 , kernel_size = 2, strides = 2 , padding = '
20        gen.add(LeakyReLU(0.3))
21
22        # fourth hidden layer
23        gen.add(Conv2D(3, kernel_size = 2 , padding = 'same', activation = 'tanh'
24
25        gen.compile(loss = 'binary_crossentropy' , optimizer = Adam(0.0001 , 0.5
26        return gen
```

```
In [35]: 1 generator = image_generator()
          2 generator.summary()
          3 visualkeras.layered_view(generator)
```

Model: "sequential\_4"


Layer (type)	Output Shape	Param #
=====		
dense_4 (Dense)	(None, 16384)	1654784
leaky_re_lu_16 (LeakyReLU)	(None, 16384)	0
reshape_3 (Reshape)	(None, 8, 8, 256)	0
conv2d_transpose_9 (Conv2DTr	(None, 16, 16, 128)	131200
leaky_re_lu_17 (LeakyReLU)	(None, 16, 16, 128)	0
conv2d_transpose_10 (Conv2DT	(None, 32, 32, 128)	65664
leaky_re_lu_18 (LeakyReLU)	(None, 32, 32, 128)	0
conv2d_transpose_11 (Conv2DT	(None, 64, 64, 128)	65664
leaky_re_lu_19 (LeakyReLU)	(None, 64, 64, 128)	0
conv2d_7 (Conv2D)	(None, 64, 64, 3)	1539
=====		
Total params: 1,918,851		
Trainable params: 1,918,851		
Non-trainable params: 0		

Out[35]:



```
In [36]: 1 noise = tf.random.normal([1, 100])
2         generated_image = generator(noise, training=False)
3
4         plt.imshow(generated_image[0, :, :, 0], cmap='gray')
```

```
Out[36]: <matplotlib.image.AxesImage at 0x20817c57bb0>
```



Above are two figures depicting the model layers and the output that we get when we give any arbitrary data to the untrained generator.

## Discriminator

The image discriminator is designed to take an input of 64 x 64 image and passes it through 3 hidden layers with the filter 128 and remaining setting as per the conv2Dtranspose. The reason for the majority of the parameters being the same is that that Conv2D and Conv2Dtranspose are simply different approaches for the same problem at hand. Flattening and drop are incorporated to guard against overfitting and to get the final output as a single value after passing through a dense layer.


```
In [37]: 1 def discriminator():
2         dis = Sequential()
3
4         # Input Layer
5         dis.add(Conv2D(64, kernel_size= 2, padding = 'same' , strides = 2 , inp
6         dis.add(layers.LeakyReLU(0.3))
7
8         # First Hidden Layer
9         dis.add(Conv2D(128 , kernel_size = 2, padding = 'same' , strides = 2 ))
10        dis.add(LeakyReLU(alpha = 0.3))
11
12        # second hidden layer
13        dis.add(Conv2D(128 , kernel_size = 2, padding = 'same' , strides = 2))
14        dis.add(LeakyReLU(alpha = 0.3))
15
16        # Third Hidden Layer
17        dis.add(Conv2D(256 , kernel_size = 2, padding = 'same' , strides = 2))
18        dis.add(LeakyReLU(alpha = 0.3))
19
20        # Flattening and drops
21        dis.add(Flatten())
22        dis.add(Dropout(0.2))
23
24        # Outplut Layer
25        dis.add(Dense(1 , activation = 'sigmoid'))
26
27        dis.compile(loss = 'binary_crossentropy', optimizer = Adam(0.0001,0.5))
28        return dis
```

```
In [38]: 1 discriminator = discriminator()
2 discriminator.summary()
3 visualkeras.layered_view(discriminator)
```

Model: "sequential\_5"

Layer (type)	Output Shape	Param #
=====		
conv2d_8 (Conv2D)	(None, 32, 32, 64)	832
leaky_re_lu_20 (LeakyReLU)	(None, 32, 32, 64)	0
conv2d_9 (Conv2D)	(None, 16, 16, 128)	32896
leaky_re_lu_21 (LeakyReLU)	(None, 16, 16, 128)	0
conv2d_10 (Conv2D)	(None, 8, 8, 128)	65664
leaky_re_lu_22 (LeakyReLU)	(None, 8, 8, 128)	0
conv2d_11 (Conv2D)	(None, 4, 4, 256)	131328
leaky_re_lu_23 (LeakyReLU)	(None, 4, 4, 256)	0
flatten_1 (Flatten)	(None, 4096)	0
dropout_1 (Dropout)	(None, 4096)	0
dense_5 (Dense)	(None, 1)	4097
=====		
Total params: 234,817		
Trainable params: 234,817		
Non-trainable params: 0		

Out[38]:



```
In [39]: 1 decision = discriminator(generated_image)
2 print (decision)
```

tf.Tensor([[0.4999248]], shape=(1, 1), dtype=float32)

After running the above line of code, we can see that we have a positive value which indicates that our discriminator is declaring the random image that we generated earlier as a valid image which shows that an untrained model is quite despicable.

## Joining the Generator and Discriminator

The following lines of code simply take in the input, reshape it to be compatible with the generator,



feeds the input to the generator to generate a fake image and subsequently feeds in the generated fake image to the discriminator. These 2 actions are then combined together to form a gan model with the loss selected to 'cross entropy' and optimizer 'Adam'.

```
In [41]: 1  ## Joining the Generator and Discriminator to make the model
          2  g_input = tf.keras.Input(shape = (dim,))
          3  fake_image = generator(g_input)
          4  g_output = discriminator(fake_image)
          5
          6  gan = tf.keras.Model(g_input , g_output)
          7  gan.compile(loss = 'binary_crossentropy', optimizer = Adam(0.0001, 0.5))
```

```
In [42]: 1  train_images_art.shape
```

```
Out[42]: (300, 64, 64, 3)
```

## Training the Model

```
In [53]: 1  batch_size = 16
          2  steps_per_epoch = 100
```

In [58]:

```

1  # Constant noise for viewing how the GAN progresses
2  static_noise = np.random.normal(0, 1, size=(100, dim))
3  save_path = "testers"
4
5  dis_loss = []
6  gen_loss = []
7
8  temp_epochs = 20
9  for epoch in range(temp_epochs):
10     for batch in range(steps_per_epoch):
11         noise = np.random.normal(0, 1, size=(batch_size, dim))
12
13         real_images = train_photo_art[np.random.randint(0, train_photo_art.s
14
15         fake_images = generator.predict(noise)
16
17         x = np.concatenate((real_images, fake_images))
18
19         disc = np.zeros(2*batch_size)
20         disc[:batch_size] = 0.9
21
22         d_loss = discriminator.train_on_batch(x, disc)
23
24         gen = np.ones(batch_size)
25         g_loss = gan.train_on_batch(noise, gen)
26
27
28
29     print(f'Epoch: {epoch} \t Discriminator Loss: {d_loss} \t\t Generator Lo
30     dis_loss.append(d_loss)
31     gen_loss.append(g_loss)
32
33     if epoch % 2 == 0:
34         generated_images = generator.predict(noise)
35         plt.figure(figsize=(10, 10))
36
37         for i, image in enumerate(generated_images):
38             plt.subplot(4, 4, i+1)
39             if channels == 1:
40                 plt.imshow(np.clip(image.reshape((64, 64)), 0.0, 1.0), cmap=
41             else:
42                 image = ((image + 1) / 2)
43                 plt.imshow(np.clip(image.reshape((64, 64, channels)), 0.0, 1
44             plt.axis('off')
45
46         plt.tight_layout()
47
48         if epoch != None:
49             plt.savefig(f'{save_path}/gan-images_epoch-{epoch}.png')

```

```

Epoch: 0          Discriminator Loss: 0.8146036863327026
rator Loss: 0.7401208281517029
Epoch: 1          Discriminator Loss: 0.5999828577041626
rator Loss: 1.094464898109436
Epoch: 2          Discriminator Loss: 0.5524779558181763
rator Loss: 1.1557050943374634
Epoch: 3          Discriminator Loss: 0.5504552125930786

```

Gene  
Gene  
Gene  
Gene

rator Loss: 1.1346198320388794	
Epoch: 4	Discriminator Loss: 0.645463228225708
ss: 1.090438723564148	
Epoch: 5	Discriminator Loss: 0.5152895450592041
rator Loss: 1.9175986051559448	
Epoch: 6	Discriminator Loss: 0.6535927653312683
rator Loss: 1.7531461715698242	
Epoch: 7	Discriminator Loss: 0.5425638556480408
rator Loss: 1.3331127166748047	
Epoch: 8	Discriminator Loss: 0.7127001285552979
rator Loss: 1.4164378643035889	
Epoch: 9	Discriminator Loss: 0.4402596354484558

Generator Lo

Gene

Gene

Gene

Gene

Gene

## Results & Analysis


```
In [59]: 1 d = {'disc_loss': dis_loss, 'gen_loss': gen_loss}
          2 loss_df = pd.DataFrame(d)
          3 loss_df
```

Out[59]:

	disc_loss	gen_loss
0	0.814604	0.740121
1	0.599983	1.094465
2	0.552478	1.155705
3	0.550455	1.134620
4	0.645463	1.090439
5	0.515290	1.917599
6	0.653593	1.753146
7	0.542564	1.333113
8	0.712700	1.416438
9	0.440260	1.570973
10	0.545479	1.455778
11	0.648317	1.084201
12	0.529102	1.502698
13	0.534861	1.236134
14	0.513120	1.265668
15	0.586240	1.202451
16	0.612298	1.366253
17	0.669679	1.092560
18	0.459349	1.503913
19	0.577431	1.131724

```
In [71]: 1 loss_df.plot.line(figsize = (15,10))
2 plt.xlabel('epochs')
3 plt.ylabel('loss')
4 plt.title('loss vs epoch')
```

```
Out[71]: Text(0.5, 1.0, 'loss vs epoch')
```



## Analysis

The graph above depicting the generator and discriminator loss shows that initially both the models are performing close to each other, however, as we increase the epochs, it seems that the discriminator is performing way better than the generator and is able to decipher quite well the images. The images themselves also have not taken a lot of logical shape but only seem to give an outline of the generated sceneries. Though the model was quite basic, it does seem to generate images which depict the shape of a valid scenery though the pixel quality seemingly lacks much clarity.

## Conclusion

The GAN is found to be a very effective tool in generating images which are not real. The most important takeaways from the project include the design of two networks which have opposing roles and design and varying the size of the images during input and during the course of passing the data through the various layers. The network does seem to be improving with the increment in

the epochs so a simple increment in the number of epochs might improve the model. Moreover, looking at the losses, the generator seems to be lacking the potential to generate good images, for that it might need some additional layers in its structure to extract the features more effectively. The monet train dataset is also quite small so techniques such as data augmentation might be quite effective as well.

## References

1. <https://rubiksgcode.net/2018/12/17/implementing-gan-dcgan-with-python/>  
(<https://rubiksgcode.net/2018/12/17/implementing-gan-dcgan-with-python/>).
2. [https://keras.io/examples/generative/dcgan\\_overriding\\_train\\_step/](https://keras.io/examples/generative/dcgan_overriding_train_step/)  
([https://keras.io/examples/generative/dcgan\\_overriding\\_train\\_step/](https://keras.io/examples/generative/dcgan_overriding_train_step/)).
3. <https://magenta.tensorflow.org/music-vae> (<https://magenta.tensorflow.org/music-vae>).
4. <https://www.tensorflow.org/tutorials/generative/dcgan>  
(<https://www.tensorflow.org/tutorials/generative/dcgan>).
5. [https://keras.io/guides/sequential\\_model/](https://keras.io/guides/sequential_model/) ([https://keras.io/guides/sequential\\_model/](https://keras.io/guides/sequential_model/)).
6. [https://keras.io/api/layers/convolution\\_layers/convolution2d/](https://keras.io/api/layers/convolution_layers/convolution2d/)  
([https://keras.io/api/layers/convolution\\_layers/convolution2d/](https://keras.io/api/layers/convolution_layers/convolution2d/)).
7. [https://keras.io/api/layers/activation\\_layers/leaky\\_relu/](https://keras.io/api/layers/activation_layers/leaky_relu/)  
([https://keras.io/api/layers/activation\\_layers/leaky\\_relu/](https://keras.io/api/layers/activation_layers/leaky_relu/)).
8. [https://keras.io/api/layers/convolution\\_layers/convolution2d\\_transpose/](https://keras.io/api/layers/convolution_layers/convolution2d_transpose/)  
([https://keras.io/api/layers/convolution\\_layers/convolution2d\\_transpose/](https://keras.io/api/layers/convolution_layers/convolution2d_transpose/)).

In [ ]:

1